

task-1

February 6, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: train_data = pd.read_csv('train_data.txt', sep=':::', names_
    ↳=['ID', 'TITLE', 'GENRE', 'DESCRIPTION'])
display(train_data.head())
train_data.shape
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_10276\61894218.py:1: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

```
train_data = pd.read_csv('train_data.txt', sep=':::', names
=['ID', 'TITLE', 'GENRE', 'DESCRIPTION'])
```

	ID	TITLE	GENRE \
0	1	Oscar et la dame rose (2009)	drama
1	2	Cupid (1997)	thriller
2	3	Young, Wild and Wonderful (1980)	adult
3	4	The Secret Sin (1915)	drama
4	5	The Unrecovered (2007)	drama

	DESCRIPTION
0	Listening in to a conversation between his do...
1	A brother and sister with a past incestuous r...
2	As the bus empties the students for their fie...
3	To help their unemployed father make ends mee...
4	The film's title refers not only to the un-re...

```
[2]: (54214, 4)
```

```
[3]: test_data = pd.read_csv('test_data.txt', sep=':::', names_
    ↳=['ID', 'TITLE', 'GENRE', 'DESCRIPTION'])
display(test_data.head())
test_data.shape
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_10276\3327045876.py:1:
ParserWarning: Falling back to the 'python' engine because the 'c' engine does
not support regex separators (separators > 1 char and different from '\s+' are
interpreted as regex); you can avoid this warning by specifying engine='python'.
```

```
test_data = pd.read_csv('test_data.txt',sep=':::', names
=['ID','TITLE','GENRE','DESCRIPTION'])
```

	ID	TITLE \
0	1	Edgar's Lunch (1998)
1	2	La guerra de papá (1977)
2	3	Off the Beaten Track (2010)
3	4	Meu Amigo Hindu (2015)
4	5	Er nu zhai (1955)

	GENRE	DESCRIPTION
0	L.R. Brane loves his life - his car, his apar...	NaN
1	Spain, March 1964: Quico is a very naughty ch...	NaN
2	One year in the life of Albin and his family ...	NaN
3	His father has died, he hasn't spoken with hi...	NaN
4	Before he was known internationally as a mart...	NaN

```
[3]: (54200, 4)
```

```
[4]: test_solution_data = pd.read_csv('test_data_solution.txt',sep=':::', names_
↳=['ID','TITLE','GENRE','DESCRIPTION'])
display(test_solution_data.head())
test_solution_data.shape
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_10276\3836370518.py:1:
ParserWarning: Falling back to the 'python' engine because the 'c' engine does
not support regex separators (separators > 1 char and different from '\s+' are
interpreted as regex); you can avoid this warning by specifying engine='python'.
```


```
test_solution_data = pd.read_csv('test_data_solution.txt',sep=':::', names
=['ID','TITLE','GENRE','DESCRIPTION'])
```

	ID	TITLE	GENRE \
0	1	Edgar's Lunch (1998)	thriller
1	2	La guerra de papá (1977)	comedy
2	3	Off the Beaten Track (2010)	documentary
3	4	Meu Amigo Hindu (2015)	drama
4	5	Er nu zhai (1955)	drama


	DESCRIPTION
0	L.R. Brane loves his life - his car, his apar...
1	Spain, March 1964: Quico is a very naughty ch...
2	One year in the life of Albin and his family ...
3	His father has died, he hasn't spoken with hi...
4	Before he was known internationally as a mart...

```
[4]: (54200, 4)
```

```
[5]: plt.figure(figsize=(15,8))
sns.countplot(x = train_data['GENRE'], order = train_data['GENRE'].
    ↳value_counts().index)
plt.xticks(rotation=50)
plt.grid(linestyle='--',color='black')
plt.title('Number of Movies per Genre')
plt.ylabel('Number of Movies')
plt.xlabel('Genre')
plt.show()
```



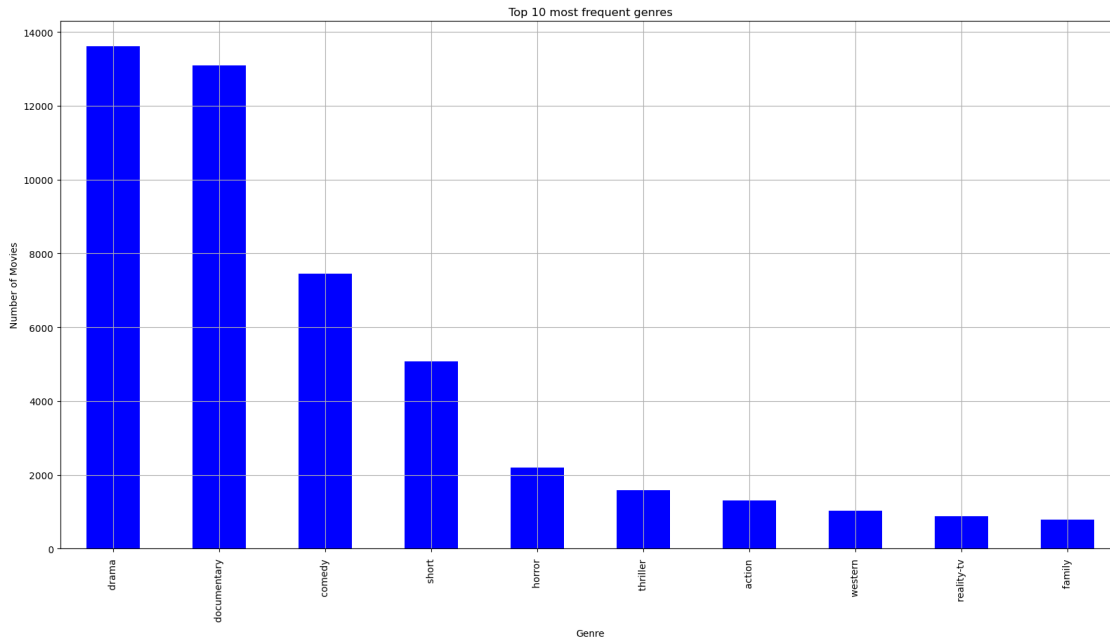
```
[6]: train_data['DESCRIPTION_length'] = train_data['DESCRIPTION'].apply(len)
plt.figure(figsize=(15,8))
sns.barplot(x='GENRE',y='DESCRIPTION_length', data=train_data)
plt.xticks(rotation=90)
plt.grid()
plt.show()
```



```
[7]: top_genre = train_data['GENRE'].value_counts().head(10)
top_genre
```

```
[7]: drama          13613
documentary       13096
comedy            7447
short             5073
horror            2204
thriller          1591
action            1315
western           1032
reality-tv         884
family             784
Name: GENRE, dtype: int64
```

```
[8]: plt.figure(figsize=(20,10))
top_genre.plot(kind='bar', color='blue')
plt.title('Top 10 most frequent genres')
plt.grid()
plt.ylabel('Number of Movies')
plt.xlabel('Genre')
plt.show()
```



```
[9]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder
from sklearn.svm import LinearSVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
```

```
[10]: train_data['DESCRIPTION'].fillna("", inplace=True)
test_data['DESCRIPTION'].fillna("", inplace=True)
```

```
[11]: train_data.isnull().sum()
```

```
[11]: ID                0
TITLE                 0
GENRE                 0
DESCRIPTION            0
DESCRIPTION_length    0
dtype: int64
```

```
[12]: test_data.isnull().sum()
```

```
[12]: ID                0
TITLE                 0
GENRE                 0
DESCRIPTION            0
dtype: int64
```

```
[13]: tvf = TfidfVectorizer(stop_words='english', max_features=100000)
x_train = tvf.fit_transform(train_data['DESCRIPTION'])
x_test = tvf.transform(test_data['DESCRIPTION'])
```

```
[14]: x_train
```

```
[14]: <54214x100000 sparse matrix of type '<class 'numpy.float64'>'
      with 2442599 stored elements in Compressed Sparse Row format>
```

```
[15]: x_test
```

```
[15]: <54200x100000 sparse matrix of type '<class 'numpy.float64'>'
      with 0 stored elements in Compressed Sparse Row format>
```

```
[16]: label = LabelEncoder()
y_train = label.fit_transform(train_data['GENRE'])
y_test = label.transform(test_solution_data['GENRE'])
```

```
[17]: y_train
```

```
[17]: array([ 8, 24,  1, ...,  7,  5, 12])
```

```
[18]: y_test
```

```
[18]: array([24,  5,  7, ...,  1,  8,  8])
```

```
[19]: x_train_sub, x_Val, y_train_sub, y_Val = train_test_split(x_train, y_train,
                                                             test_size=0.2,
                                                             random_state=111)
print(x_train_sub.shape)
print(x_Val.shape)
print(y_train_sub.shape)
print(y_Val.shape)
```

```
(43371, 100000)
```

```
(10843, 100000)
```

```
(43371,)
```

```
(10843,)
```

```
[20]: clf = LinearSVC()
      clf.fit(x_train_sub, y_train_sub)

      y_val_predict = clf.predict(x_Val)
```

```
C:\Users\Admin\anaconda3\Lib\site-packages\sklearn\svm\_classes.py:32:
FutureWarning: The default value of `dual` will change from `True` to `auto`
in 1.5. Set the value of `dual` explicitly to suppress the warning.
  warnings.warn(
```

```
[21]: print('Validation Accuracy:', accuracy_score(y_Val, y_val_predict))
      print('validation classification report : \n',
            ↪ classification_report(y_Val, y_val_predict))
```

Validation Accuracy: 0.5819422669003044

validation classification report :

	precision	recall	f1-score	support
0	0.43	0.32	0.36	258
1	0.74	0.39	0.51	127
2	0.42	0.19	0.26	162
3	0.65	0.17	0.27	101
4	0.00	0.00	0.00	54
5	0.53	0.57	0.55	1499
6	0.15	0.03	0.05	108
7	0.69	0.84	0.76	2649
8	0.55	0.71	0.62	2693
9	0.39	0.11	0.18	149
10	0.07	0.02	0.03	54
11	0.93	0.66	0.77	38
12	0.00	0.00	0.00	50
13	0.62	0.65	0.63	463
14	0.63	0.46	0.53	138
15	0.00	0.00	0.00	48
16	0.20	0.01	0.03	70
17	0.38	0.15	0.21	34
18	0.51	0.26	0.34	169
19	0.22	0.04	0.07	128
20	0.51	0.30	0.37	132
21	0.43	0.35	0.38	993
22	0.63	0.42	0.51	73
23	0.56	0.29	0.39	78
24	0.34	0.20	0.25	324
25	0.67	0.06	0.11	32
26	0.84	0.84	0.84	219
accuracy			0.58	10843
macro avg	0.45	0.30	0.33	10843
weighted avg	0.55	0.58	0.55	10843

C:\Users\Admin\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\Admin\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Admin\anaconda3\Lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
[22]: y_pred = clf.predict(x_test)
print('test Accuracy:', accuracy_score(y_test, y_pred))
print('test classification report : \n', classification_report(y_test, y_test))
```

test Accuracy: 0.09357933579335793

test classification report :

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1314
1	1.00	1.00	1.00	590
2	1.00	1.00	1.00	775
3	1.00	1.00	1.00	498
4	1.00	1.00	1.00	264
5	1.00	1.00	1.00	7446
6	1.00	1.00	1.00	505
7	1.00	1.00	1.00	13096
8	1.00	1.00	1.00	13612
9	1.00	1.00	1.00	783
10	1.00	1.00	1.00	322
11	1.00	1.00	1.00	193
12	1.00	1.00	1.00	243
13	1.00	1.00	1.00	2204
14	1.00	1.00	1.00	731
15	1.00	1.00	1.00	276
16	1.00	1.00	1.00	318
17	1.00	1.00	1.00	181
18	1.00	1.00	1.00	883
19	1.00	1.00	1.00	672
20	1.00	1.00	1.00	646
21	1.00	1.00	1.00	5072
22	1.00	1.00	1.00	431
23	1.00	1.00	1.00	391
24	1.00	1.00	1.00	1590
25	1.00	1.00	1.00	132
26	1.00	1.00	1.00	1032
accuracy			1.00	54200
macro avg	1.00	1.00	1.00	54200
weighted avg	1.00	1.00	1.00	54200


```
[23]: from sklearn.naive_bayes import MultinomialNB
      mnb = MultinomialNB()
      mnb.fit(x_train,y_train)
```

```
[23]: MultinomialNB()
```

```
[24]: mnb.predict(x_test)
```

```
[24]: array([8, 8, 8, ..., 8, 8, 8])
```

```
[25]: from sklearn.linear_model import LogisticRegression
      lr= LogisticRegression()
      lr.fit(x_train, y_train)
```

```
C:\Users\Admin\anaconda3\Lib\site-
packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[25]: LogisticRegression()
```

```
[26]: y_pred = lr.predict(x_test)
```

```
[27]: print('Accuracy score:', accuracy_score(y_test,y_pred))
```

```
Accuracy score: 0.2511439114391144
```

```
[ ]:
```